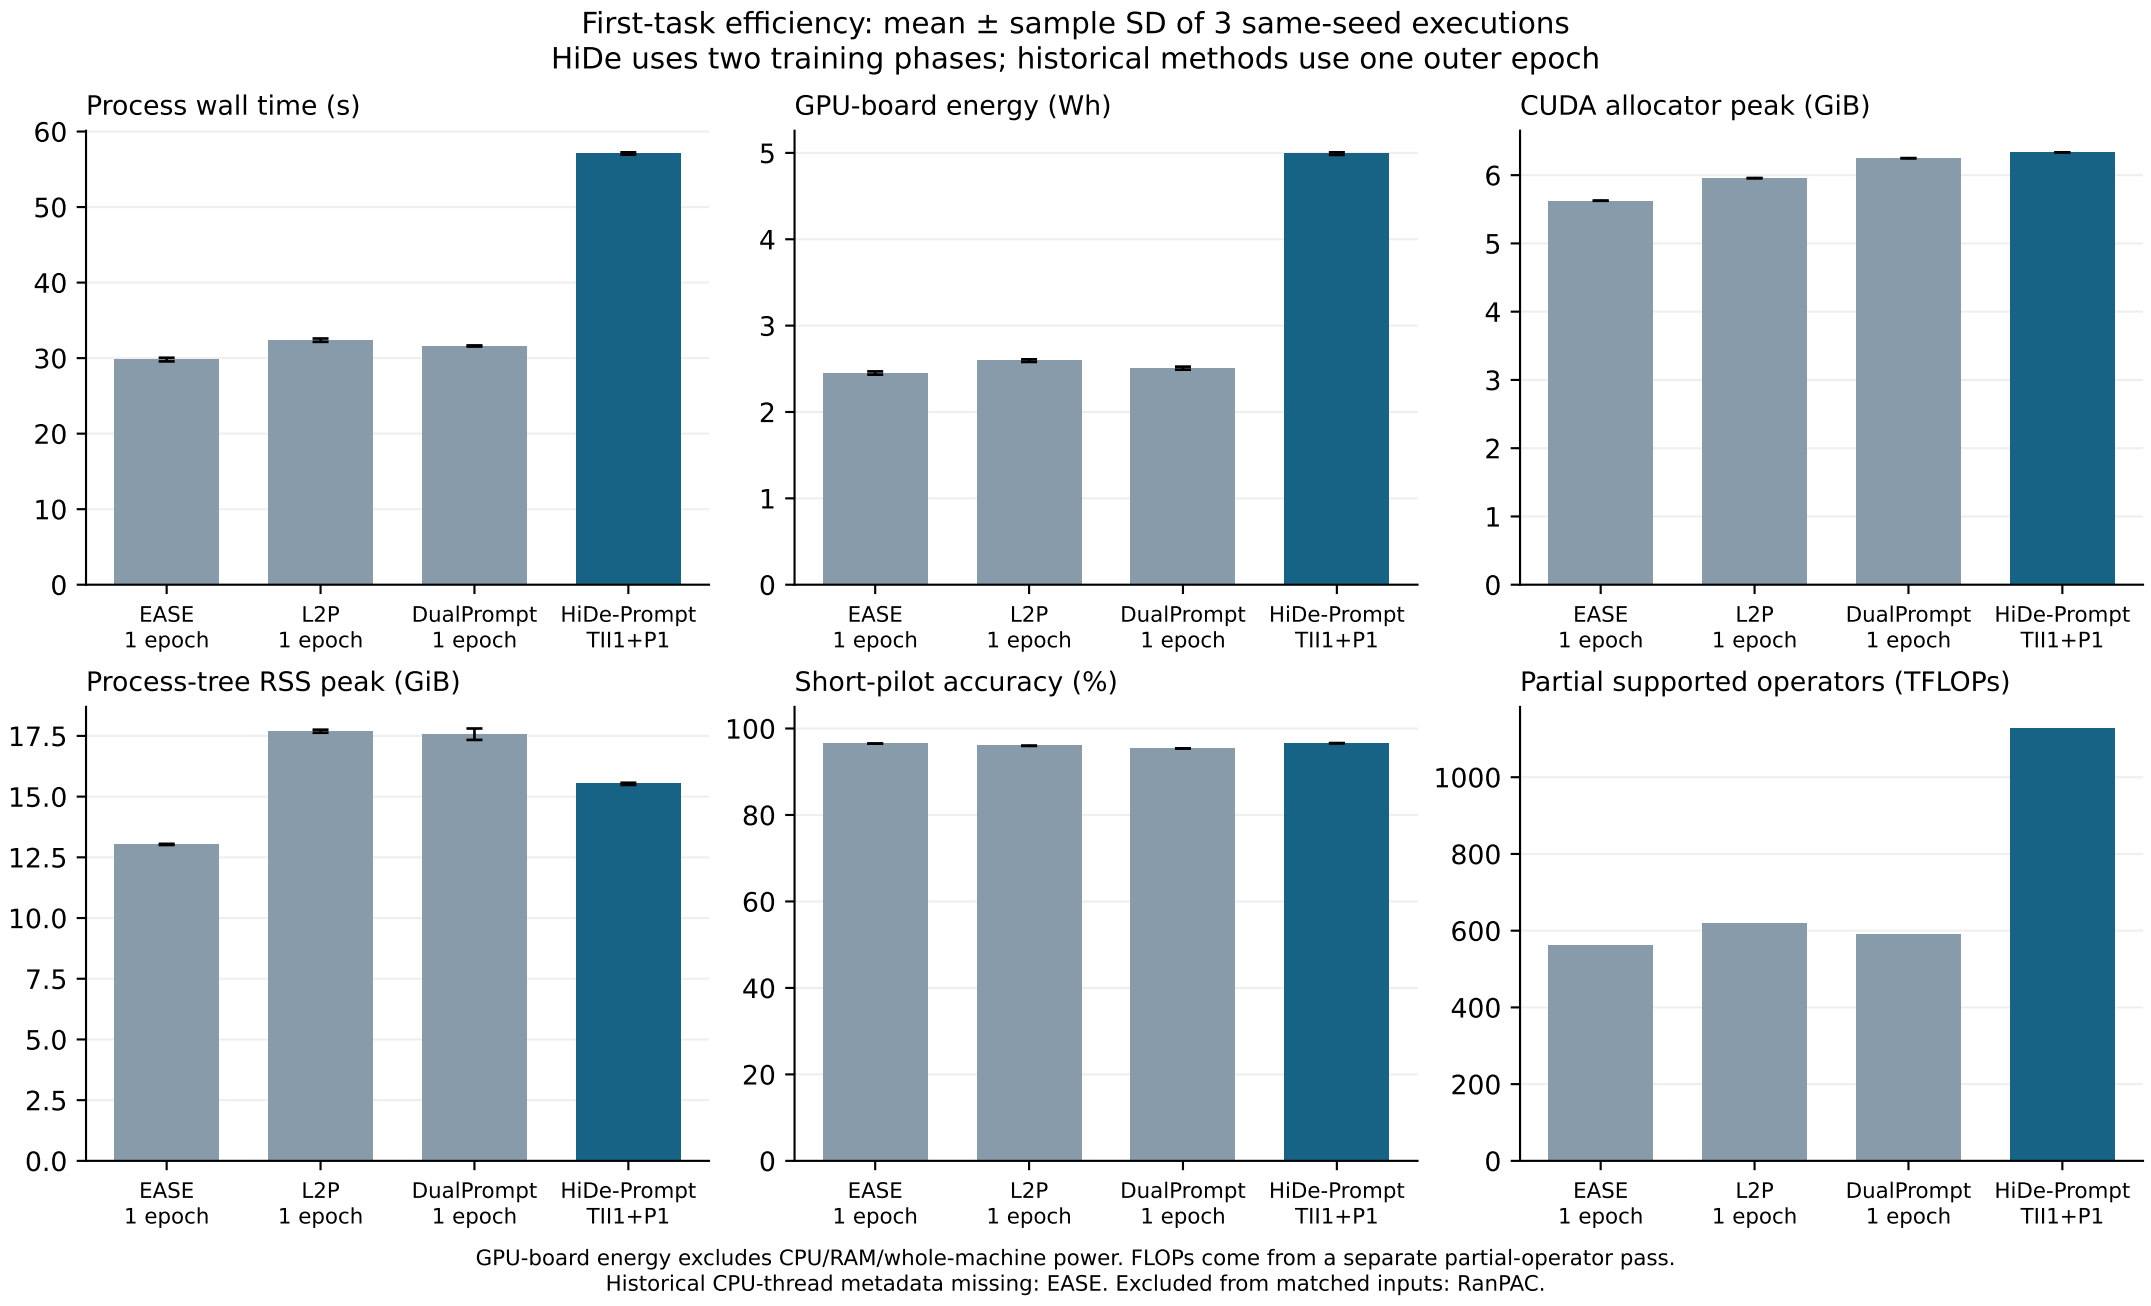
First-task efficiency: mean ± sample SD of 3 same-seed executions
HiDe uses two training phases; historical methods use one outer epoch
Process wall time (s)
GPU-board energy (Wh)
CUDA allocator peak (GiB)
Process-tree RSS peak (GiB)
Short-pilot accuracy (%)
Partial supported operators (TFLOPs)
EASE 1 epoch
L2P 1 epoch
DualPrompt 1 epoch
HiDe-Prompt TII1+P1
GPU-board energy excludes CPU/RAM/whole-machine power. FLOPs come from a separate partial-operator pass.
Historical CPU-thread metadata missing: EASE. Excluded from matched inputs: RanPAC.